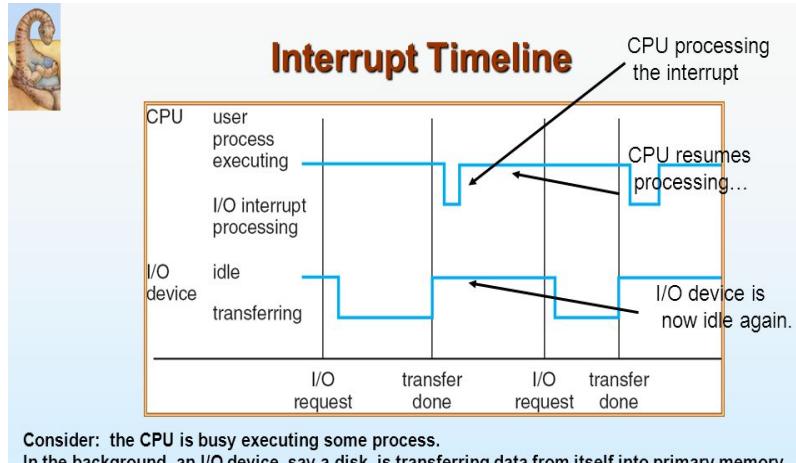
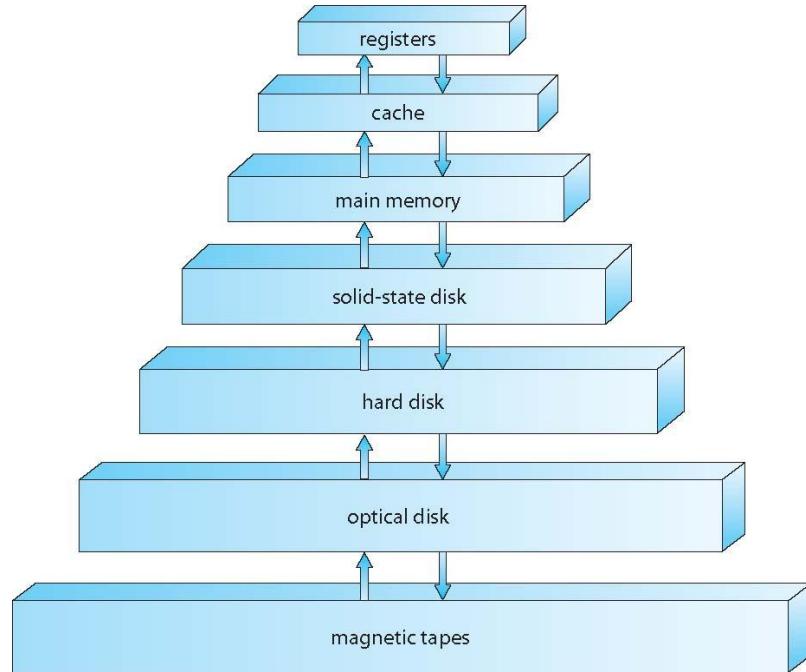


A computer system can be divided roughly into four components: the hardware, the operating system, the application programs, and the users (as shown in figure) Abstract view of the components of a computer system.

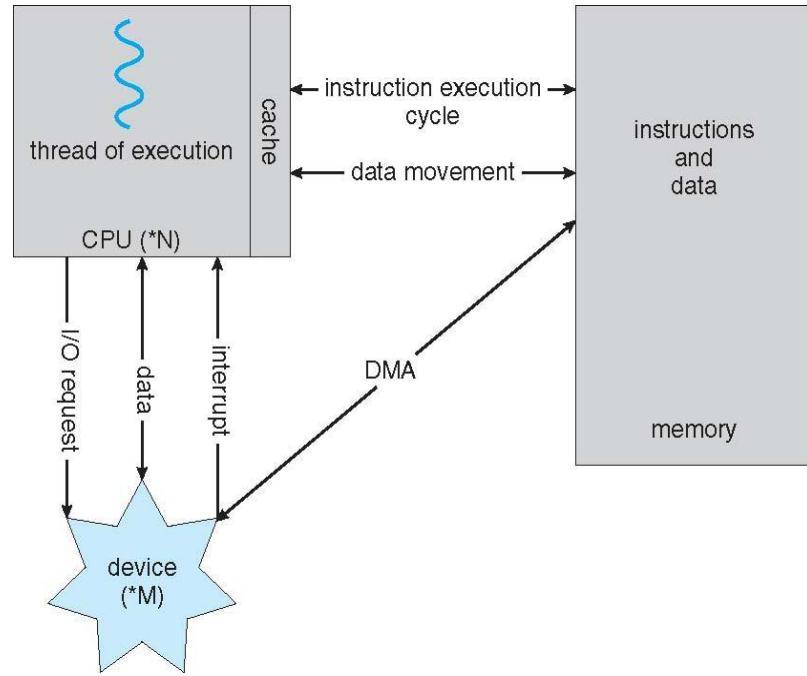


Consider: the CPU is busy executing some process.
 In the background, an I/O device, say a disk, is transferring data from itself into primary memory.
 When the transfer is complete, a signal (interrupt) is sent to the CPU.
 The CPU can suspend 'current' operations, and process the I/O interrupt
 Of course at this time, the I/O device becomes idle
 Note the CPU time to process the interrupt is very small (much smaller than chart indicates).
 Then the CPU resumes normal processing (whatever that might be)
 It might resume suspended process or start executing a different process.

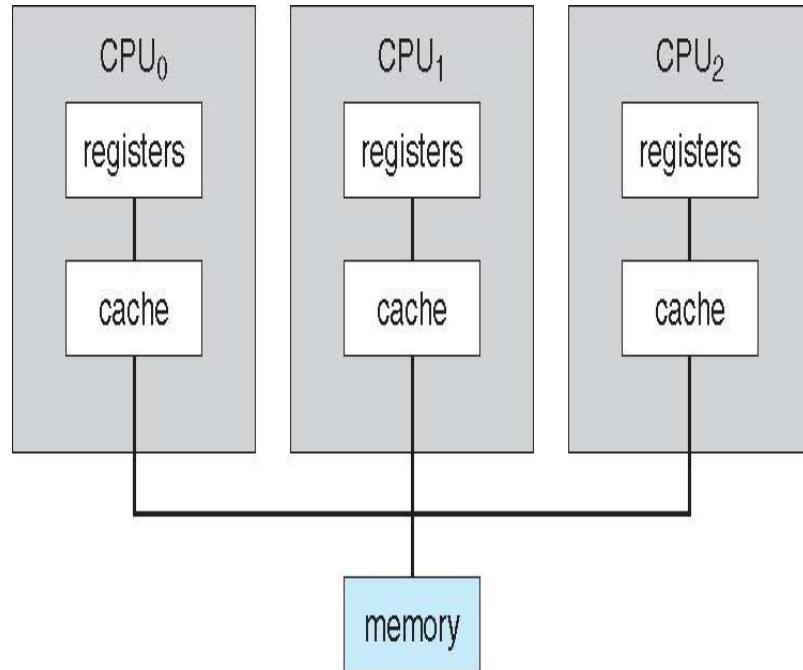
This is sufficient for now. Much more later.



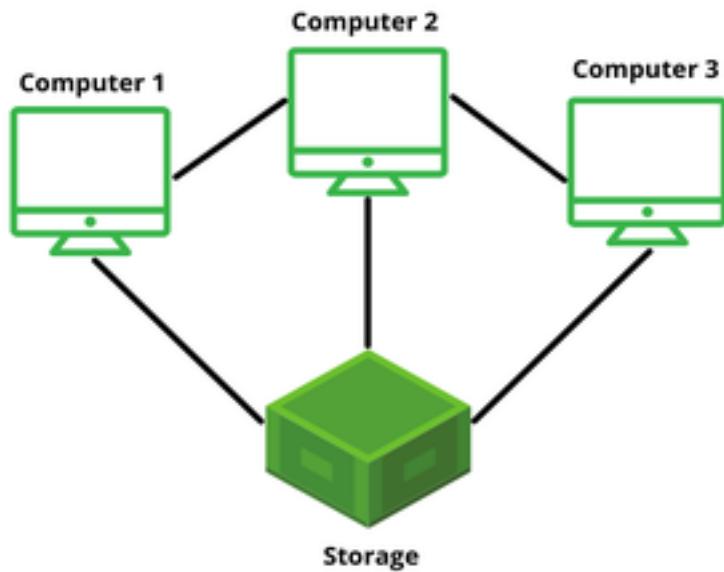
Storage-device hierarchy



Working of I/O Operation

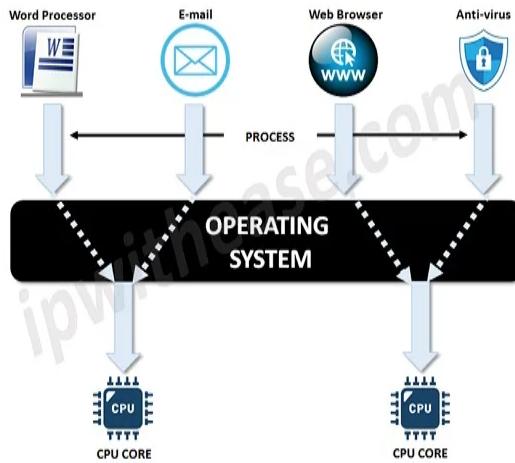


The most common systems use symmetric multiprocessing (SMP), in which each processor performs all tasks within the operating system. SMP means that all processors are peers; no master-slave relationship exists between processors. Figure illustrates a typical SMP architecture. Symmetric multiprocessing architecture.



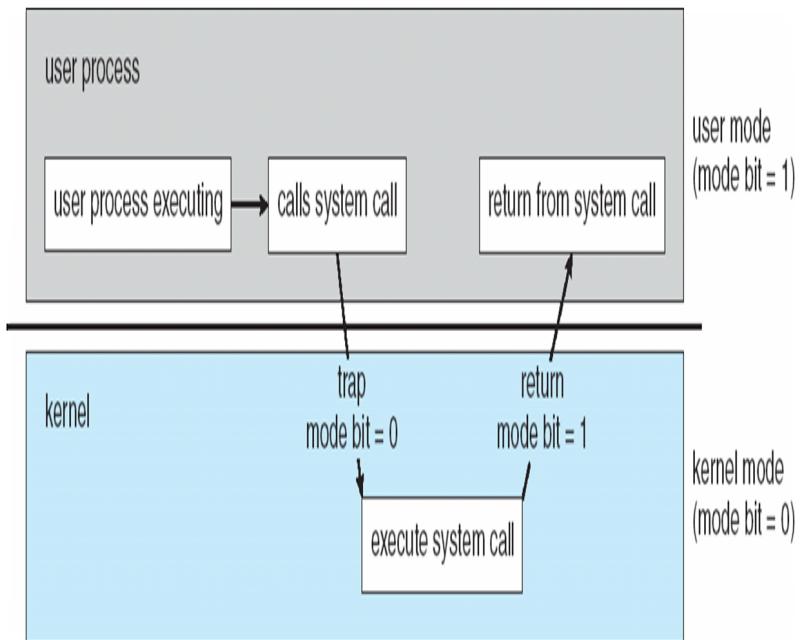
Another type of multiple CPU system is the clustered system. Like multiprocessor systems, clustered systems gather together multiple CPUs to accomplish computational work. Clustered systems differ from multiprocessor systems, however, in that they are composed of two or more individual systems coupled together called as cluster. Provides high availability.

Multi-Tasking

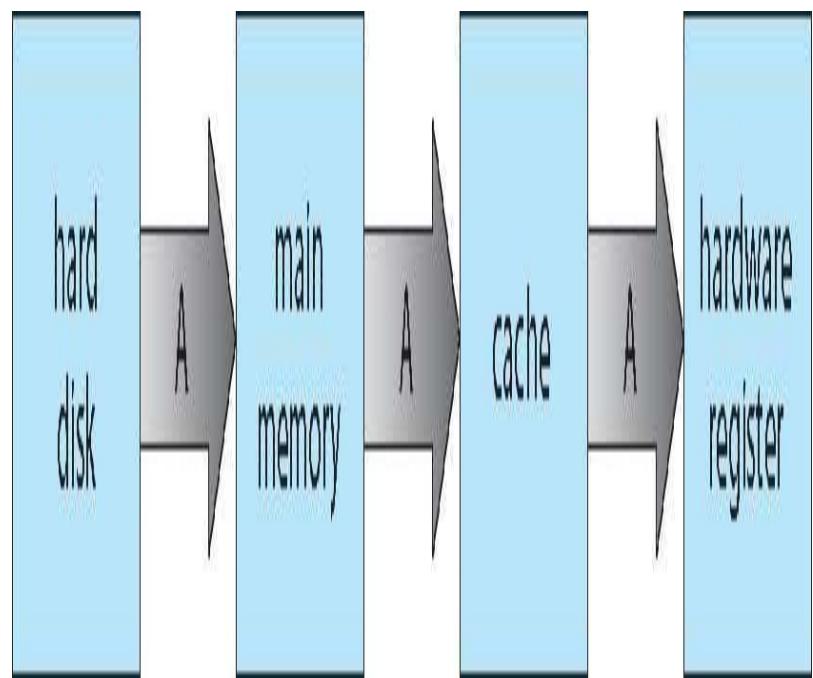


networkinterview.com
(An Initiative By ipwitthease.com)

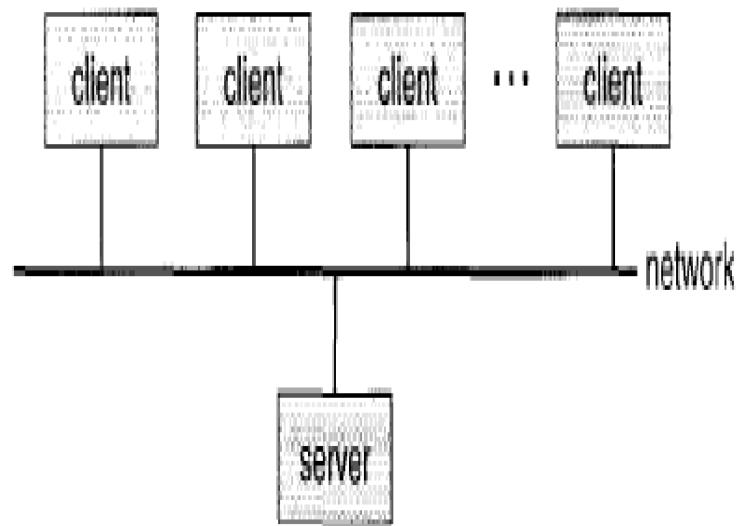
Dept. of ISE



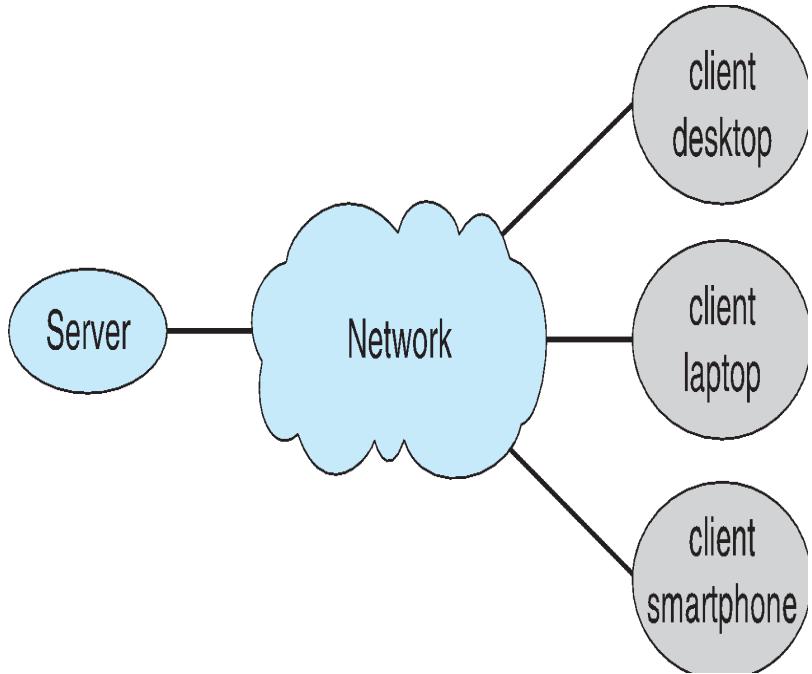
Dept. of ISE



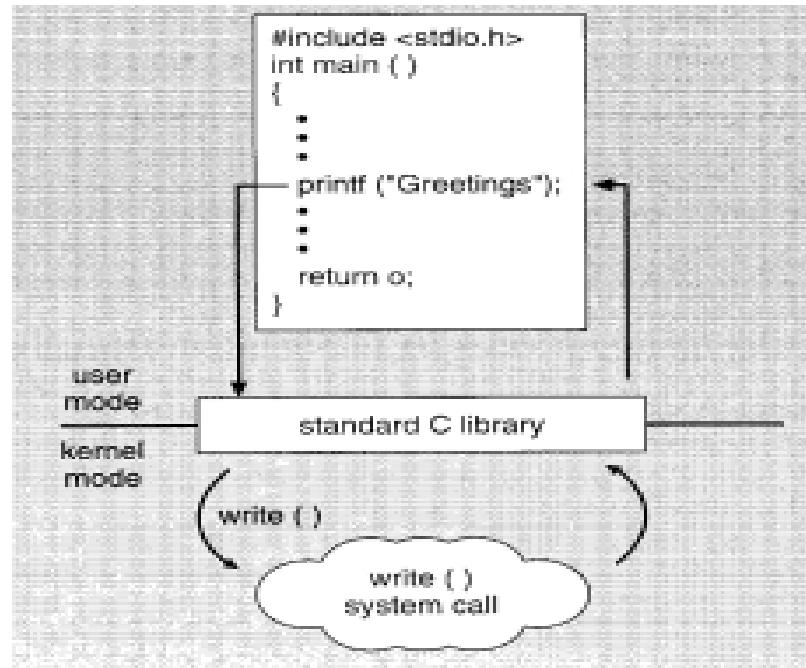
Dept. of ISE



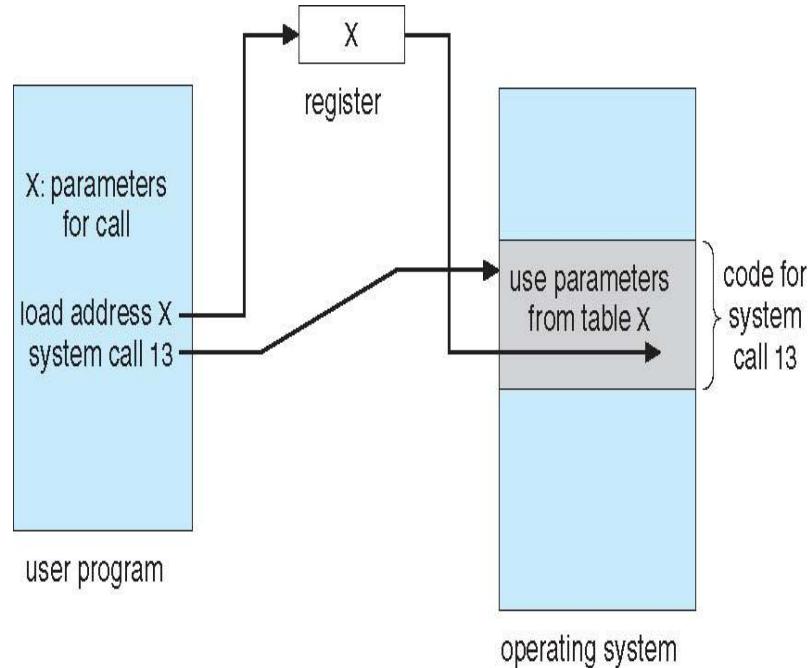
Dept. of ISE



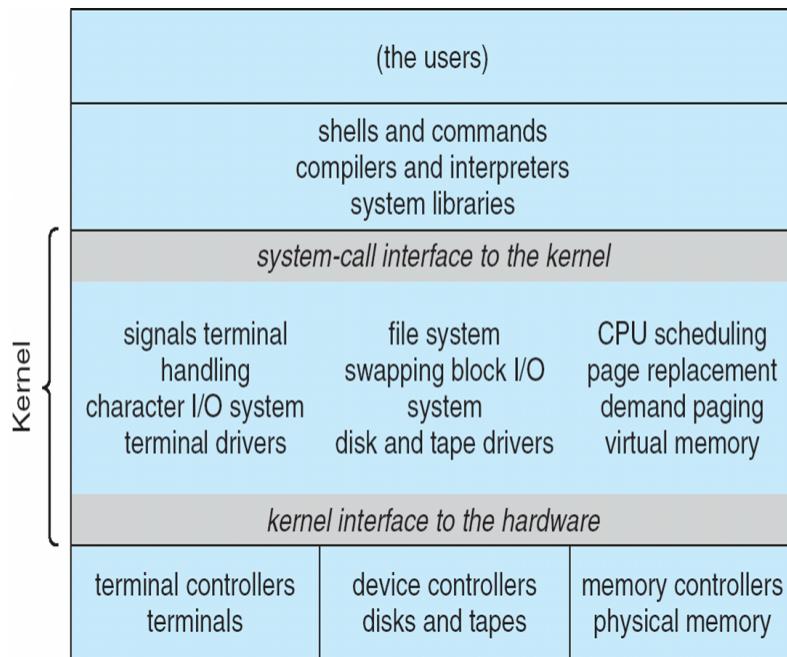
Dept. of ISE



The relationship between an API, the system-call interface, and the operating system is shown in Figure 1.13, which illustrates how the operating system handles a user application invoking the `open()` system call.



Dept. of ISE



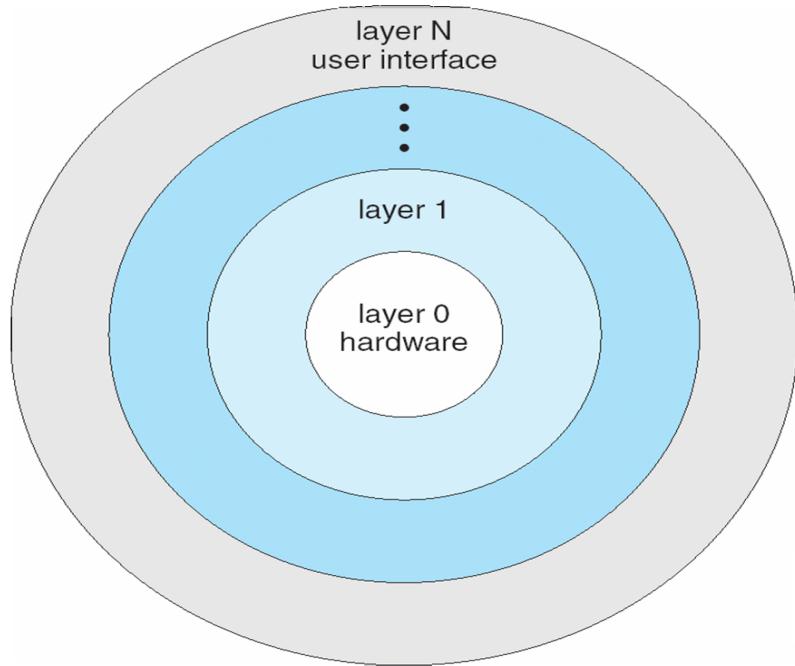
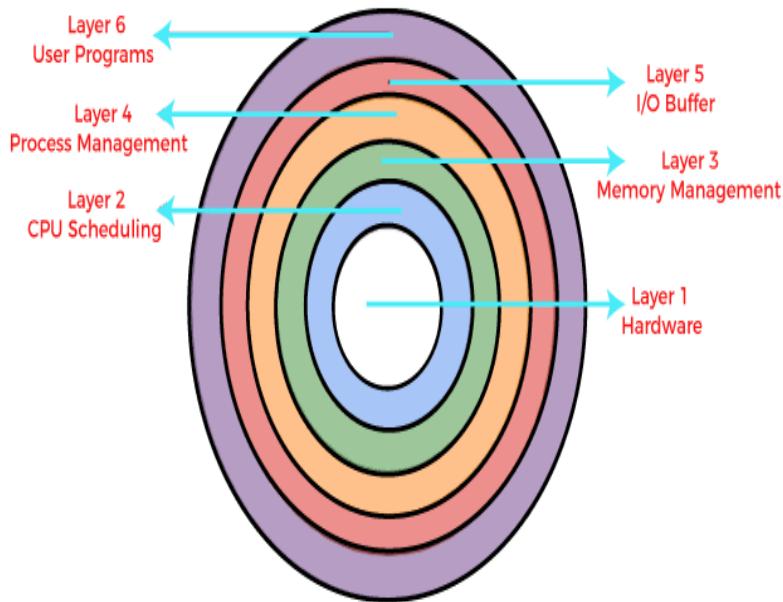
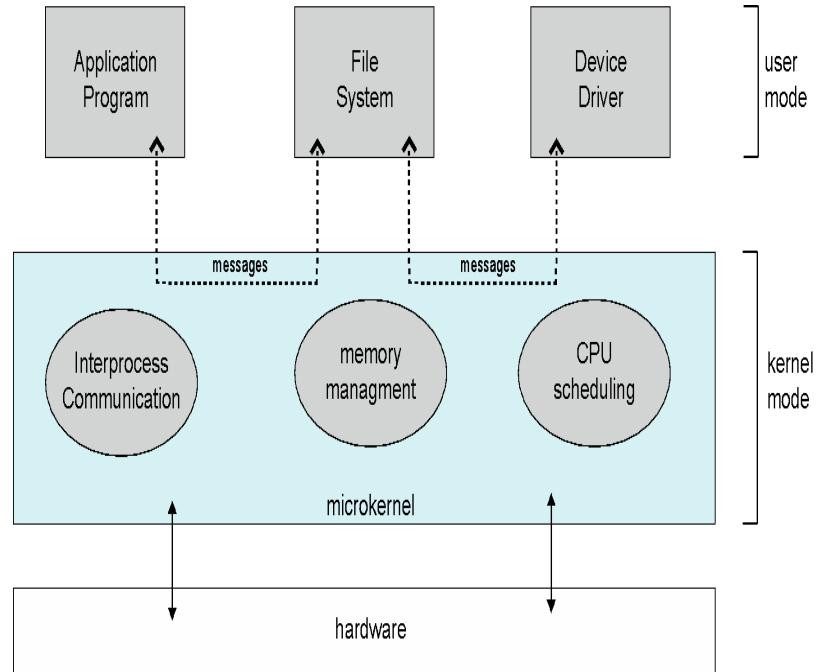


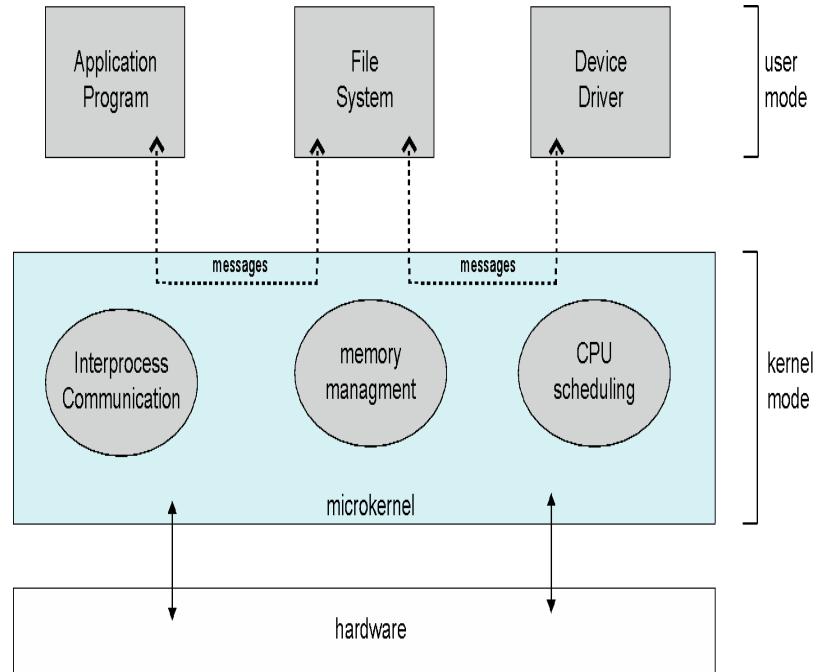
Figure: A layered operating system



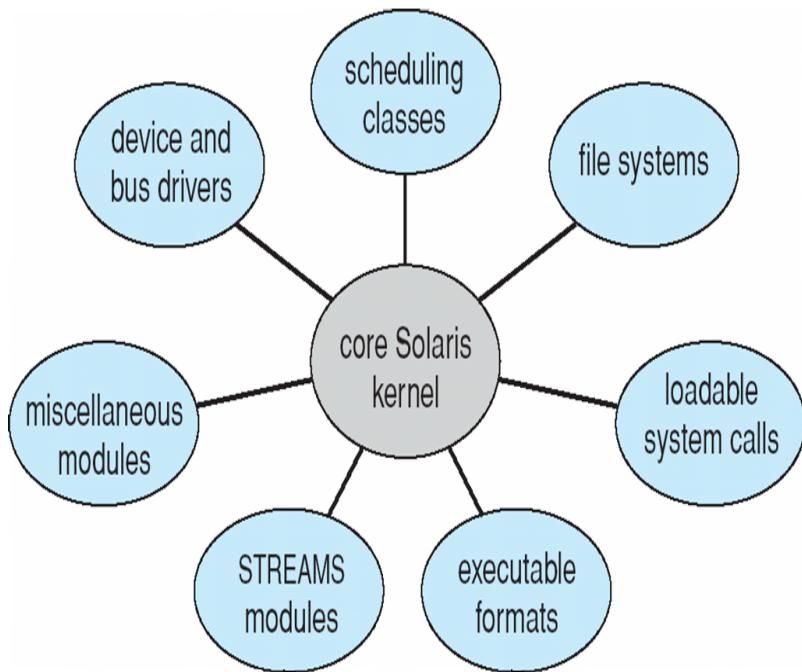
A final problem with layered implementations is that they tend to be less efficient than other types. For instance, when a user program executes an I/O operation, it executes a system call that is trapped to the I/O layer, which calls the memory-management layer, which in turn calls the CPU-scheduling layer, which is then passed to the hardware. At each layer, the parameters may be modified, dat



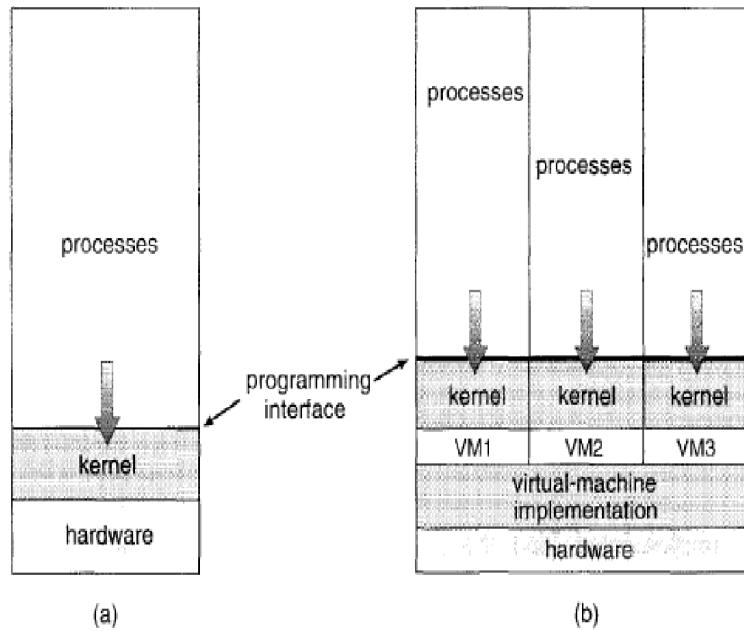
Dept. of ISE



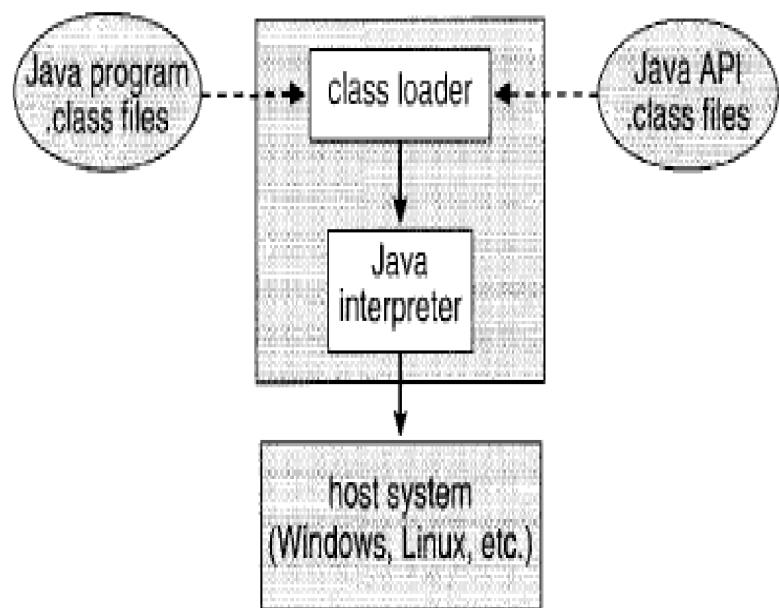
Dept. of ISE



Dept. of ISE



Dept. of ISE



Dept. of ISE

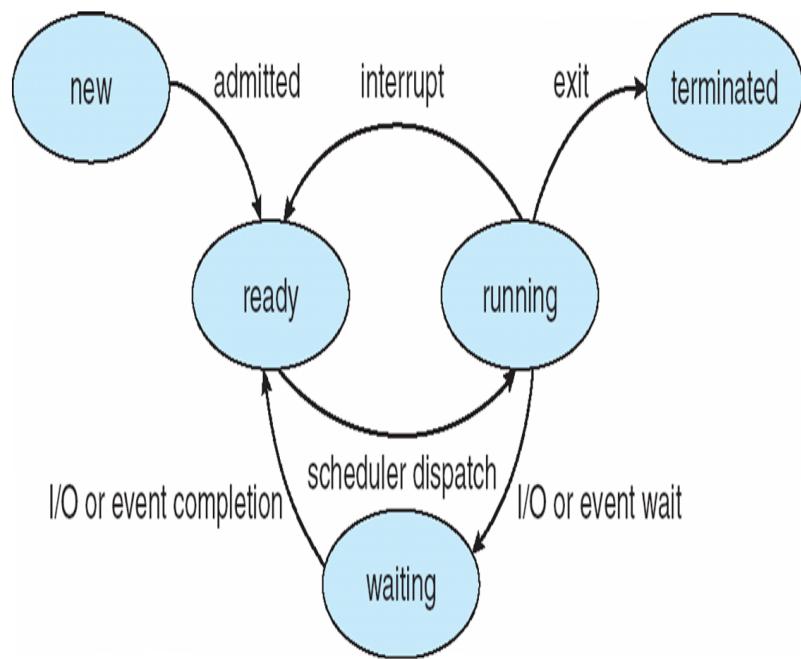
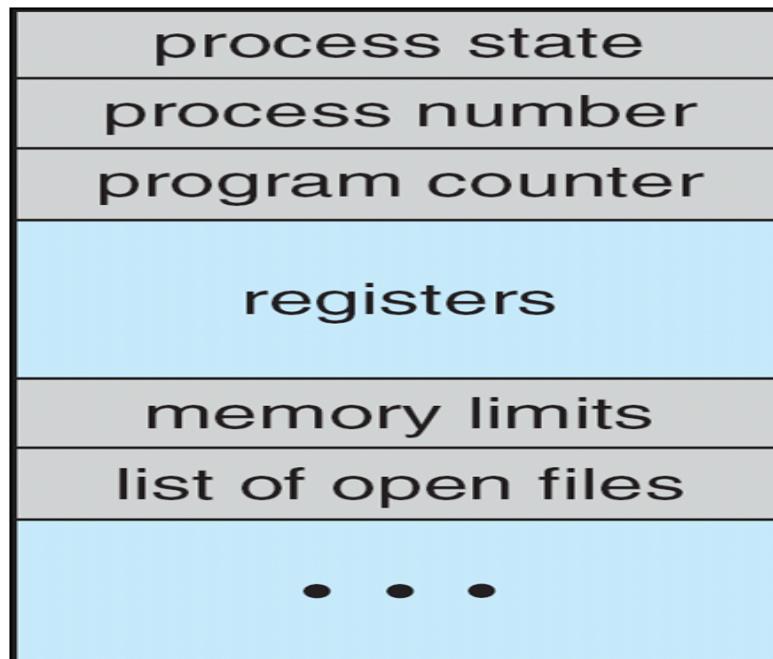
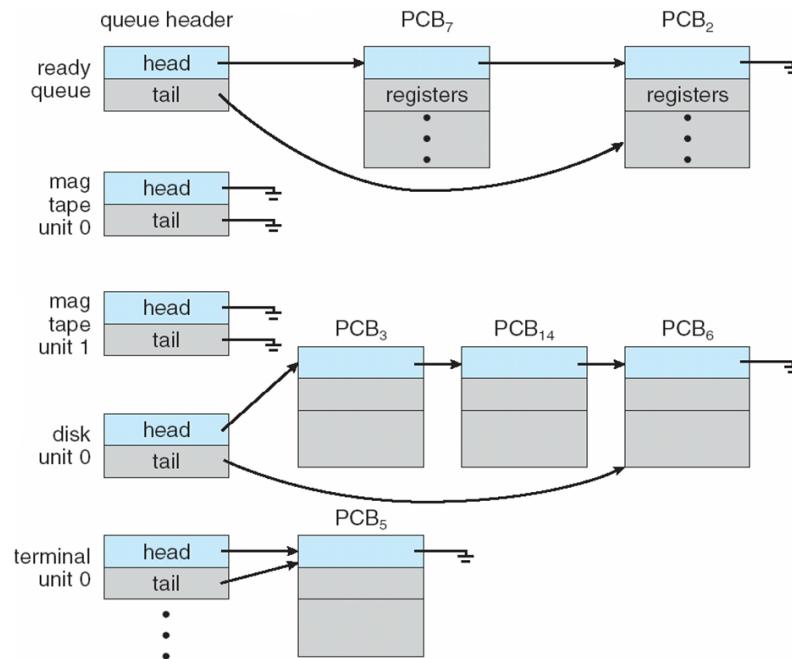


Diagram of process states.



Process Control Block Each process is represented in the operating system by a process control block (PCB) also called a task control block. A PCB is shown in Figure. It contains many pieces of information associated with a specific process, including these:



Dept. of ISE

```

#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

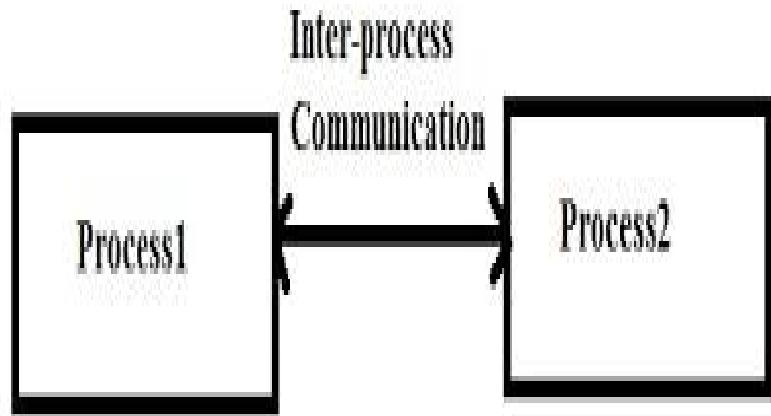
int main()
{
pid_t pid;

/* fork a child process */
pid = fork();

if (pid < 0) /* error occurred */
    fprintf(stderr, "Fork Failed");
    exit(-1);
}
else if (pid == 0) /* child process */
    execvp("/bin/ls","ls",NULL);
}
else /* parent process */
/* parent will wait for the child to complete */
wait(NULL);
printf("Child Complete");
exit(0);
}
}

```

To illustrate these differences, let's first consider the UNIX operating system. In UNIX, as we've seen, each process is identified by its process identifier,



©Elprocus.com

Dept. of ISE